

Using Futhark for a fast, parallel implementation of  
the Simultaneous Iterative Reconstruction Technique -  
A pre-study

Mette Bjerg and Lærke Pedersen

October 29, 2018

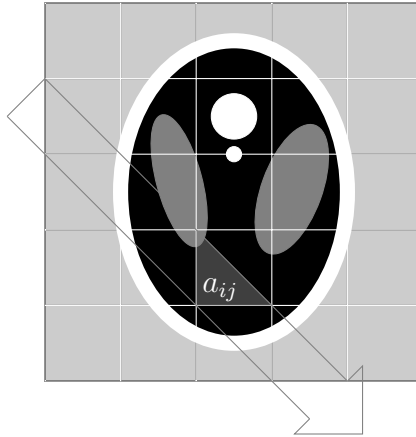


Figure 1: Example of a weighting  $a_{ij}$ .

## 1 Introduction

Computed tomography is the inverse problem of reconstructing an image or volume from its x-ray projections. The x-ray source spins around the object to be analyzed and sends x-rays that hit a detector on the opposite side. The detector shows how much of each x-ray was attenuated when passing through the volume. In this report we will focus on a type of algebraic reconstruction algorithm called the Simultaneous Iterative Reconstruction Algorithm (SIRT) and investigate how we may use the high level data-parallel language Futhark for implementing a fast version of this algorithm. The reason for choosing this algorithm is that it provides good reconstruction quality under non-optimal circumstances, but unfortunately has very poor performance compared to techniques such as filtered back-projection.

In algebraic reconstruction we solve the problem as a linear system of equations. The main idea is that the process can be modelled as a linear transformation by discretizing the object to be reconstructed into  $N$  pixels. First we place a coordinate system with origo at the center of the object to be reconstructed and denote by  $\boldsymbol{\theta}$  the vector of angles between the positions of the source and origo. For each angle several x-rays are cast from the source. We denote by the vector  $\boldsymbol{\rho}$  the signed distances from each line to origo. The data produced by the process is called the sinogram. Then the sinogram values  $p_i$  for each  $(\theta_k, \rho_l)$  are a weighted sum of the attenuations at each pixel  $f_j$  that the  $(\theta_k, \rho_l)$  ray passes through:

$$\sum_{j=1}^N a_{ij} f_j = p_i \quad (1)$$

Where  $a_{ij}$  are the weights, corresponding to the fraction of the pixel  $j$  that the ray  $i$  covers.

Writing all the projections as a column vector  $\mathbf{p}$  and the attenuation values to be reconstructed as a column vector  $\mathbf{f}$  the weightings are represented as an  $M \times N$  matrix

$\mathbf{A}$ , we obtain a linear system of the form:

$$\mathbf{p} = \mathbf{A}\mathbf{f} \quad (2)$$

These systems of equations may easily be solved under the right circumstances, where  $M = N$ . However this is rarely the case. In most real cases  $M > N$ , i.e. the number of projections is larger than the number of pixels to be reconstructed and the size of the matrix is very large - more about this in the next section.

However, the algebraic reconstruction methods also have some advantages. Since the model closely relates to the real world scenario the weightings can be refactored to take irregularities in the setup, such as differences in beam energies or irregular geometries and missing data into account. Furthermore these methods generally give better image quality than analytic methods when the data is sparse.

A linear system like this is typically solved by minimizing some norm:

$$\|\mathbf{A}\mathbf{f} - \mathbf{p}\| \quad (3)$$

An example is the SIRT algorithm. The action of the matrix  $\mathbf{A}$  is called the *forward projection*, and the matrix itself is called the *system matrix*, *projection matrix* or *radon matrix*. Each row of  $\mathbf{A}$  represents the coefficients of the equation for one ray. The action of the transposed projection matrix  $\mathbf{A}^T$  on the sinogram vector is called the *backprojection*, and can be visualized as smearing the projection values across the reconstruction in the direction of the ray, or equivalently summing up the contributions of each ray for a given pixel. The idea behind the SIRT algorithm is to forward project the current reconstruction, then subtract this from the original projection data and do a weighted backprojection resulting in a correction factor which can be added to the current reconstruction. The update equation is:

$$\mathbf{f}^n = \mathbf{f}^{(n-1)} + \mathbf{C}\mathbf{A}^T\mathbf{R}(\mathbf{p} - \mathbf{A}\mathbf{f}^{(n-1)}), \quad (4)$$

where  $\mathbf{C}$  and  $\mathbf{R}$  are the diagonal matrices containing the inverse column and row sums of the system matrix respectively.

It can be shown that this iterative scheme solves the problem:

$$\mathbf{f}^* = \operatorname{argmin}_{\mathbf{f}} \|\mathbf{p} - \mathbf{A}\mathbf{f}\|_{\mathbf{R}}, \quad (5)$$

where  $\|\mathbf{x}\|_{\mathbf{R}} = \mathbf{x}^T\mathbf{R}\mathbf{x}$ .

The backprojection and forward projection operations are standard operations in many iterative algebraic reconstruction methods and are the bottle necks of the algorithms [1].

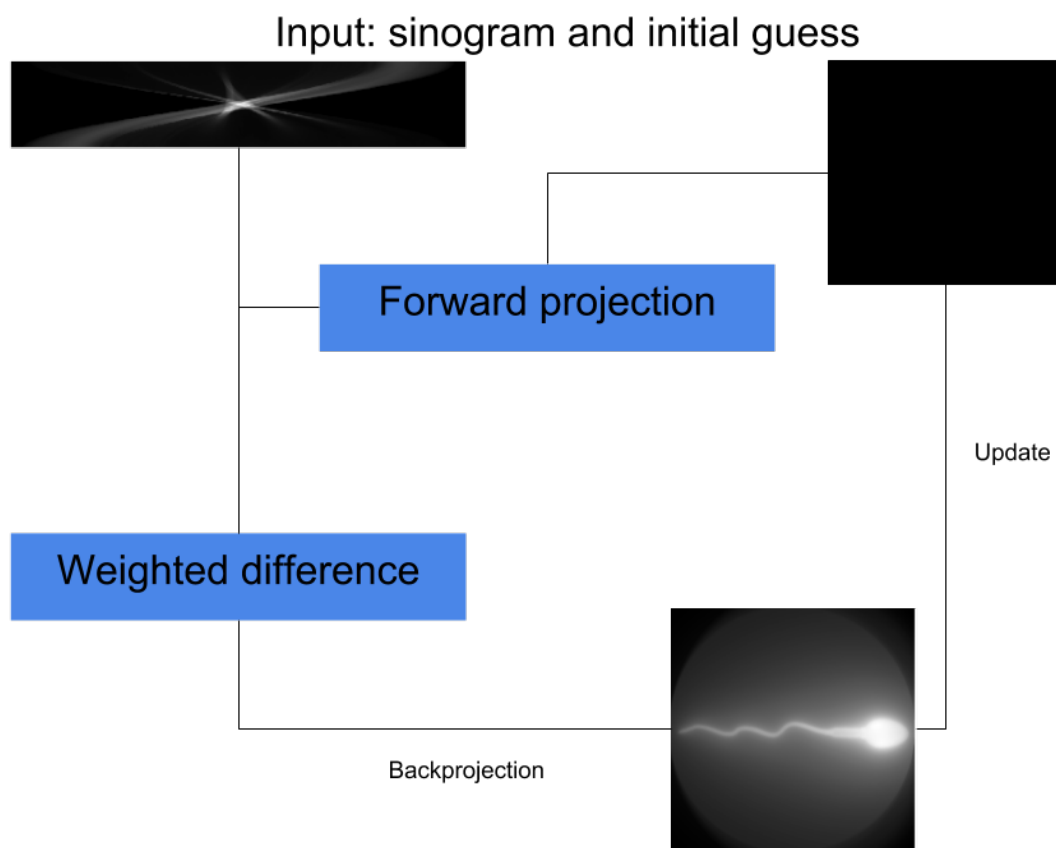


Figure 2: An illustration of the SIRT algorithm.

Therefore, our main focus has been on optimizing these operations, and then combining these to the SIRT algorithm.

Different beam geometries exist, such as parallel beams, fan beams and cone beams. The geometry must be considered when constructing the system matrix. However as this study is intended for applications with synchrotron data only parallel beam geometries will be considered. Furthermore, we will only consider reconstructions of 2D images, since this is simply a pre-study and it makes the problem conceptually easier to follow.

## 2 Problem size

One of the main problems when parallelizing the algorithm is that the amount of data in real applications is huge. Images from synchrotrons are generated with detectors of sizes up to  $4000 \times 4000$ . To accurately generate 3D reconstructions it has been proven that approximately  $\frac{\pi \cdot N}{2}$  viewing angles are needed, where  $N$  is the size of the detector in one direction [2].

We will benchmark the algorithms by using sizes  $N$  ranging from 128 to 4096 and use  $\lceil \frac{\pi \cdot N}{2} \rceil$  angles and  $N$  lines for each of these sizes.

To solve the large problems it is not possible to store the whole system matrix on the GPU. Lets take as an example the amount of space used to store the matrix for a 3D problem with a  $N = 4000$  detector, with  $\lceil \frac{\pi \cdot N}{2} \rceil$  angles, and 4000 rays per angle. The matrix is sparse having at most  $2 \cdot N - 1$  entries in each row, so instead of storing the whole matrix with all the zeroes, we will consider a semi-sparse representation with one 2 dimensional array containing the data as floats in arrays on each row of size  $2 \cdot N - 1$ , and one matching 2 dimensional array containing the column indexes of the datapoints. For the example sketched above the number of rows of each matrix will be  $\lceil \frac{\pi \cdot 4000}{2} \rceil \cdot 4000$  and the number of columns is  $2 \cdot 4000 - 1$ . Assuming ints and floats both take up 4 bytes each then storing the matrix requires  $4 \cdot 2 \cdot \lceil \frac{\pi \cdot 4000}{2} \rceil \cdot 4000 \cdot (2 \cdot 4000 - 1) \approx 1TB$ . This is of course one of the largest problems, but even for a problem with  $N = 512$  we're looking at the matrix taking up around  $3 - 4GB$ . Besides the matrix we also need to store the sinogram data which will take approximately  $420MB$ , and the reconstruction taking up around  $270MB$ . Hence even for relatively small problems we're looking at an estimated  $5GB$  of data. GPU memory is expensive, and is typically in the range 2-8GB for consumer end cards, and 10-48GB for high end cards. Thus, the small problems are not able to fit in a standard consumer card, and large problems won't fit on any cards.

There are several ways of getting around this problem. Some problems exhibit symmetries in the system matrix so that we only have to store part of the matrix. This is already somewhat accounted for in the above calculations. In the example we only considered the matrix for one slice of the reconstruction since it will be equivalent for other slices when using parallel beam geometry. Thus for cone beam geometries, the memory needed would be approximately  $N$  times greater. Other symetries involve angle symmetries, where for example a scan from a  $45^\circ$  or  $135^\circ$  angle would contain the same values for mirrored pix-

Lærke  
should  
probably  
check  
these  
calcula-  
tions...

els. However these are obviously not generally applicable since they will depend on which angles are used in the reconstruction.

Another more flexible approach is to calculate the matrix values on the fly which is what we have done. For this we used the code from a bachelor project in which they did an implementation of finding the system matrix using futhark. We had two different parallel implementations in futhark and a sequential version in python which we compared. We decided to use the fastest version XXX

### 3 Nested parallel version

A first approach was to do a nested parallel version of forward and back projection using on the fly (OTF) computation of the system matrix.

A looped version of the forward projection with the system matrix cut in *steps* chunks looks like this:

```

1 for step = 0; step < steps; step++
2   A = getRays(raysperstep)
3   for ray = 0; ray < raysperstep; ray++
4     acc = 0.0
5     for p = 0; p<numpixels; p++
6       acc+= A[ray][p]*image[p]
7   FP[step*raysperstep+ray] = acc

```

Figure 3: A looped version of the forward projection, where the rowsperstep should be the largest number possible such that the computations fit in the memory.  $step * rowsperstep$  should equal the total number of rows.

A looped version of the backprojections looks like this:

```

1 for step = 0; step < steps; step++
2   A = getRays(raysperstep)
3   AT = A.transpose()
4   for p = 0; p<numpixels; p++
5     acc = 0.0
6     for ray = 0; ray<raysperstep; ray++
7       acc+= AT[p][ray]*sinogram[ray]
8   BP[p] += acc

```

Figure 4: A looped version of the back projection, where the rowsperstep should be the largest number possible such that the computations fit in the memory.  $step * rowsperstep$  should equal the total number of rows.

## 4 Flattening

Exploiting nested parallelism, as in the code sketched in the previous section, is difficult on GPU since the hardware is organized on one or two parallel levels that allow threads to communicate via shared scratchpad memory.

One way to get around this problem is to use a flattening transformation. The problem with this is that it will require even more memory usage, and may prevent opportunities for locality optimizations .

## 5 Comparison to a CUDA implementation

We compared our implementation of forward projection and back projection to implementations from the astra toolbox.

## 6 Memory coalescence and other optimization ideas

## References

- [1] Y. Long, J. A. Fessler, and J. M. Balter. 3d forward and back-projection for x-ray ct using separable footprints. *IEEE Transactions on Medical Imaging*, 29(11):1839–1850, Nov 2010.
- [2] F. Natterer. *The Mathematics of Computerized Tomography*. Society for Industrial and Applied Mathematics, 2001.
- [3] Z. Xue, L. Zhang, and J. Pan. A new algorithm for calculating the radiological path in ct image reconstruction. In *Proceedings of 2011 International Conference on Electronic Mechanical Engineering and Information Technology*, volume 9, pages 4527–4530, Aug 2011.